

Hsiang Yu (Anna) Huang

617-319-5044 | huanganna1004@gmail.com | linkedin.com/in/hsiangyuhuang | hsiangyuhuang.com

AI Engineer / New Grad (Applied LLM / Backend)

Product-oriented AI engineer with hands-on experience building production LLM systems (RAG, agentic workflows, memory) and scalable backend services. Proven ability to improve model reliability, deploy containerized services on AWS, and support real user-facing products in startup and research environments.

EDUCATION

Boston University

M.S. in Data Science (GPA: 3.6)

Boston, MA

Jan. 2026

- **Award:** Winner – DS+X Hackathon 2025 (Best Overall for *RhettSearch* AI tool)
- **Coursework:** Deep Learning, Data Engineering, Graduate Databases, Artificial Intelligence

National Taiwan University of Science and Technology

B.B.A. in Industrial Management & Finance; Minor in Computer Science (GPA: 3.8)

Taipei, Taiwan

Jun. 2023

TECHNICAL SKILLS

AI / LLM Systems: RAG, Agentic Workflows (LangGraph, LangChain), Prompt Engineering, Tool Calling, Vector Databases (Pinecone, MongoDB Atlas), Context & Memory Management, LLM Evaluation (LangSmith)

Backend & Cloud: Python (FastAPI), TypeScript (Node.js, Next.js), PostgreSQL, MongoDB, Redis, Kafka, Docker, CI/CD, AWS (EC2, S3, CloudFront), Linux, Supabase

Data Engineering: ETL Pipelines, PySpark, Azure Synapse, PostHog, Power BI, Looker Studio

EXPERIENCE

AI Engineer Intern

Finz (FinTech Startup)

Dec. 2025 – Jan. 2026

Boston, MA

- Engineered a production-grade ETL pipeline using **FastAPI** and **Kafka** to normalize financial data from Stripe and Quickbook, enabling real-time cashflow analysis for AI-driven insights.
- Improved LLM intent routing accuracy by 20%, reducing fallback responses and increasing successful query resolution in a customer-facing financial assistant.
- Orchestrated containerized backend services using **Docker** and **GitHub Actions** on AWS EC2, ensuring deployment consistency across environments.

Research Assistant – LLM Platform

BU BIT Lab

Sep. 2025 – Jan. 2026

Boston, MA

- Architected a context-aware memory system using **LangGraph** that combines short-term conversational history with vector-based retrieval to prevent context drift in long-running AI sessions.
- Developed an agentic query routing framework to classify user intent and orchestrate specialized prompt chains, integrating external tool-calling for real-time data grounding.
- Designed a high-granularity MongoDB logging schema to track latency, token usage, and interaction signals, facilitating systematic LLM evaluation and performance tuning.

Full Stack Developer

Citale (BU Spark! Launch Lab)

Sep. 2024 – May 2025

Boston, MA

- Developed and launched core social features including real-time messaging and event-based maps using **Next.js** and **Supabase**, supporting successful beta testing with local businesses.
- Optimized PostgreSQL database performance through normalized schema design and strict indexing, ensuring data integrity for scalable user interactions.

PROJECTS

Multi-Agent LLM Narrative System (Long-Context Generation) | *LangGraph, Supabase*

- Designed a hierarchical multi-agent system (Director, Planner, Writer) for long-form narrative generation, overcoming context window limits via structured planning and iterative generation.
- Implemented dynamic memory summarization to maintain long-term plot coherence while reducing token consumption by 30% during long-chapter generation.

Relational Database Engine Kernel | *C++20*

- Built a storage engine kernel featuring a buffer pool manager with O(1) LRU eviction and a B+ Tree index supporting efficient search and cascading splits.